

```

import pandas as pd
import numpy as np
import seaborn as sns
from scipy import stats
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn import metrics
from sklearn.metrics import accuracy_score
from sklearn.metrics import roc_auc_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.metrics import roc_curve
from sklearn.metrics import precision_recall_curve
from sklearn.model_selection import train_test_split, KFold,
cross_val_score
from sklearn.preprocessing import MinMaxScaler

data = pd.read_csv('logistic_regression.csv')
data.columns

Index(['loan_amnt', 'term', 'int_rate', 'installment', 'grade',
      'sub_grade',
      'emp_title', 'emp_length', 'home_ownership', 'annual_inc',
      'verification_status', 'issue_d', 'loan_status', 'purpose',
      'title',
      'dti', 'earliest_cr_line', 'open_acc', 'pub_rec', 'revol_bal',
      'revol_util', 'total_acc', 'initial_list_status',
      'application_type',
      'mort_acc', 'pub_rec_bankruptcies', 'address'],
      dtype='object')

data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 27 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   loan_amnt                            396030 non-null  float64
1   term                                396030 non-null  object
2   int_rate                             396030 non-null  float64
3   installment                          396030 non-null  float64
4   grade                                396030 non-null  object
5   sub_grade                           396030 non-null  object
6   emp_title                           373103 non-null  object
7   emp_length                          377729 non-null  object
8   home_ownership                      396030 non-null  object
9   annual_inc                          396030 non-null  float64
10  verification_status                 396030 non-null  object
11  issue_d                             396030 non-null  object

```

12	loan_status	396030	non-null	object
13	purpose	396030	non-null	object
14	title	394275	non-null	object
15	dti	396030	non-null	float64
16	earliest_cr_line	396030	non-null	object
17	open_acc	396030	non-null	float64
18	pub_rec	396030	non-null	float64
19	revol_bal	396030	non-null	float64
20	revol_util	395754	non-null	float64
21	total_acc	396030	non-null	float64
22	initial_list_status	396030	non-null	object
23	application_type	396030	non-null	object
24	mort_acc	358235	non-null	float64
25	pub_rec_bankruptcies	395495	non-null	float64
26	address	396030	non-null	object

dtypes: float64(12), object(15)

memory usage: 81.6+ MB

Mapping of target variable -

```
data['loan_status'] = data.loan_status.map({'Fully Paid':0, 'Charged Off':1})
```

```
data.isnull().sum()/len(data)*100
```

loan_amnt	0.000000
term	0.000000
int_rate	0.000000
installment	0.000000
grade	0.000000
sub_grade	0.000000
emp_title	5.789208
emp_length	4.621115
home_ownership	0.000000
annual_inc	0.000000
verification_status	0.000000
issue_d	0.000000
loan_status	0.000000
purpose	0.000000
title	0.443148
dti	0.000000
earliest_cr_line	0.000000
open_acc	0.000000
pub_rec	0.000000
revol_bal	0.000000
revol_util	0.069692
total_acc	0.000000
initial_list_status	0.000000
application_type	0.000000
mort_acc	9.543469
pub_rec_bankruptcies	0.135091

```
address          0.000000
dtype: float64
```

```
# Checking the distribution of outcome labels -
data.loan_status.value_counts(normalize=True)*100
```

```
0    80.387092
1    19.612908
Name: loan_status, dtype: float64
```

```
#the data seems imbalanced we will be needing to oversample it later
```

```
# Statistical summary of the dataset -
data.describe(include='all')
```

	loan_amnt	term	int_rate	installment
grade \				
count	396030.000000	396030	396030.000000	396030.000000
396030				
unique	NaN	2	NaN	NaN
7				
top	NaN	36 months	NaN	NaN
B				
freq	NaN	302005	NaN	NaN
116018				
mean	14113.888089	NaN	13.639400	431.849698
NaN				
std	8357.441341	NaN	4.472157	250.727790
NaN				
min	500.000000	NaN	5.320000	16.080000
NaN				
25%	8000.000000	NaN	10.490000	250.330000
NaN				
50%	12000.000000	NaN	13.330000	375.430000
NaN				
75%	20000.000000	NaN	16.490000	567.300000
NaN				
max	40000.000000	NaN	30.990000	1533.810000
NaN				

	sub_grade	emp_title	emp_length	home_ownership
annual_inc ... \				
count	396030	373103	377729	396030
3.960300e+05 ...				
unique	35	173105	11	6
NaN ...				
top	B3	Teacher	10+ years	MORTGAGE
NaN ...				
freq	26655	4389	126041	198348
NaN ...				

mean	NaN	NaN	NaN	NaN
7.420318e+04	...			
std	NaN	NaN	NaN	NaN
6.163762e+04	...			
min	NaN	NaN	NaN	NaN
0.000000e+00	...			
25%	NaN	NaN	NaN	NaN
4.500000e+04	...			
50%	NaN	NaN	NaN	NaN
6.400000e+04	...			
75%	NaN	NaN	NaN	NaN
9.000000e+04	...			
max	NaN	NaN	NaN	NaN
8.706582e+06	...			

	open_acc	pub_rec	revol_bal	revol_util \
count	396030.000000	396030.000000	3.960300e+05	395754.000000
unique	NaN	NaN	NaN	NaN
top	NaN	NaN	NaN	NaN
freq	NaN	NaN	NaN	NaN
mean	11.311153	0.178191	1.584454e+04	53.791749
std	5.137649	0.530671	2.059184e+04	24.452193
min	0.000000	0.000000	0.000000e+00	0.000000
25%	8.000000	0.000000	6.025000e+03	35.800000
50%	10.000000	0.000000	1.118100e+04	54.800000
75%	14.000000	0.000000	1.962000e+04	72.900000
max	90.000000	86.000000	1.743266e+06	892.300000


	total_acc	initial_list_status	application_type
mort_acc \			
count	396030.000000	396030	396030
358235.000000			
unique	NaN	2	3
NaN			
top	NaN	f	INDIVIDUAL
NaN			
freq	NaN	238066	395319
NaN			
mean	25.414744	NaN	NaN
1.813991			
std	11.886991	NaN	NaN
2.147930			
min	2.000000	NaN	NaN
0.000000			
25%	17.000000	NaN	NaN
0.000000			
50%	24.000000	NaN	NaN
1.000000			
75%	32.000000	NaN	NaN

```
3.000000
max      151.000000      NaN      NaN
34.000000
```

	pub_rec_bankruptcies	address
count	395495.000000	396030
unique	NaN	393700
top	NaN	USS Smith\r\nFPO AP 70466
freq	NaN	8
mean	0.121648	NaN
std	0.356174	NaN
min	0.000000	NaN
25%	0.000000	NaN
50%	0.000000	NaN
75%	0.000000	NaN
max	8.000000	NaN

```
[11 rows x 27 columns]
```


```
plt.figure(figsize=(12, 8))
sns.heatmap(data.corr(method='spearman'), annot=True, cmap='viridis')
plt.show()
```



#Comment about the correlation between Loan Amount and Installment features.

We noticed almost perfect correlation between "loan_amnt" the "installment" feature.
installment: The monthly payment owed by the borrower if the loan originates. loan_amnt: The listed amount of the loan applied for by the borrower. If at some point in time, the credit department reduces the loan amount, then it will be reflected in this value. So, we can drop either one of those columns.

```
data.drop(columns=['installment'], axis=1, inplace=True)
plt.figure(figsize=(12, 8))
sns.heatmap(data.corr(method='spearman'), annot=True, cmap='inferno')
plt.show()
```



```
data.loan_status.value_counts()
```

```
0    318357
```

```
1     77673
```

```
Name: loan_status, dtype: int64
```

Data Exploration -

1. The no of people those who have fully paid are 318357 and that of Charged Off are 77673.

```
data.groupby(by='loan_status')['loan_amnt'].describe()
```

	count	mean	std	min	25%
50% \ loan_status					
0	318357.0	13866.878771	8302.319699	500.0	7500.0
12000.0					
1	77673.0	15126.300967	8505.090557	1000.0	8525.0
14000.0					

	75%	max
loan_status		
0	19225.0	40000.0
1	20000.0	40000.0

#2. The majority of people have home ownership as Mortgage and Rent
data['home_ownership'].value_counts()

MORTGAGE	198348
RENT	159790
OWN	37746
OTHER	112
NONE	31
ANY	3

Name: home_ownership, dtype: int64

#Combining the minority classes as 'OTHER'.
data.loc[(data.home_ownership == 'ANY') | (data.home_ownership ==
'NONE'), 'home_ownership'] = 'OTHER'
data.home_ownership.value_counts()

MORTGAGE	198348
RENT	159790
OWN	37746
OTHER	146

Name: home_ownership, dtype: int64

Checking the distribution of 'Other' -
data.loc[data['home_ownership']=='OTHER',
'loan_status'].value_counts()

0	123
1	23

Name: loan_status, dtype: int64

data['issue_d'] = pd.to_datetime(data['issue_d'])
data['earliest_cr_line'] = pd.to_datetime(data['earliest_cr_line'])

*#Saw some issues in title (Looks like it was filled manually and needs
some fixing).*

data['title'].value_counts()[:20]

Debt consolidation	152472
Credit card refinancing	51487
Home improvement	15264
Other	12930
Debt Consolidation	11608
Major purchase	4769
Consolidation	3852
debt consolidation	3547
Business	2949

Debt Consolidation Loan	2864
Medical expenses	2742
Car financing	2139
Credit Card Consolidation	1775
Vacation	1717
Moving and relocation	1689
consolidation	1595
Personal Loan	1591
Consolidation Loan	1299
Home Improvement	1268
Home buying	1183

Name: title, dtype: int64

```
data['title'] = data.title.str.lower()
data.title.value_counts()[:10]
```



debt consolidation	168108
credit card refinancing	51781
home improvement	17117
other	12993
consolidation	5583
major purchase	4998
debt consolidation loan	3513
business	3017
medical expenses	2820
credit card consolidation	2638

Name: title, dtype: int64




Visualization - The grade of majority of people those who have fully paid the loan is 'B' and have subgrade 'B3'. So from where we can infer that people with grade 'B' and subgrade 'B3' are more likely to fully pay the loan.

People with grades 'A' are more likely to fully pay their loan. (T/F) so this is false

```
plt.figure(figsize=(15, 10))
plt.subplot(2, 2, 1)
grade = sorted(data.grade.unique().tolist())
sns.countplot(x='grade', data=data, hue='loan_status', order=grade)
plt.subplot(2, 2, 2)
sub_grade = sorted(data.sub_grade.unique().tolist())
g = sns.countplot(x='sub_grade', data=data, hue='loan_status',
order=sub_grade)
g.set_xticklabels(g.get_xticklabels(), rotation=90);
```



```
plt.figure(figsize=(15, 20))
plt.subplot(4, 2, 1)
sns.countplot(x='term', data=data, hue='loan_status')
plt.subplot(4, 2, 2)
sns.countplot(x='home_ownership', data=data, hue='loan_status')
plt.subplot(4, 2, 3)
sns.countplot(x='verification_status', data=data, hue='loan_status')
plt.subplot(4, 2, 4)
g = sns.countplot(x='purpose', data=data, hue='loan_status')
g.set_xticklabels(g.get_xticklabels(), rotation=90);
```




```

#Name the top 2 afforded job titles.
#Manager and Teacher are the most afforded loan job titles

plt.figure(figsize=(15, 12))
plt.subplot(2, 2, 1)
order = ['< 1 year', '1 year', '2 years', '3 years', '4 years', '5 years',
        '6 years', '7 years', '8 years', '9 years', '10+ years',]
g = sns.countplot(x='emp_length', data=data, hue='loan_status',
order=order)
g.set_xticklabels(g.get_xticklabels(), rotation=90);
plt.subplot(2, 2, 2)
plt.barh(data.emp_title.value_counts()[:30].index,
data.emp_title.value_counts()[:30])
plt.title("The most 30 jobs title afforded a loan")
plt.tight_layout()

```



#Feature Engineering -

```

def pub_rec(number):
    if number == 0.0:
        return 0
    else:
        return 1

def mort_acc(number):
    if number == 0.0:
        return 0
    elif number >= 1.0:
        return 1
    else:
        return number

```

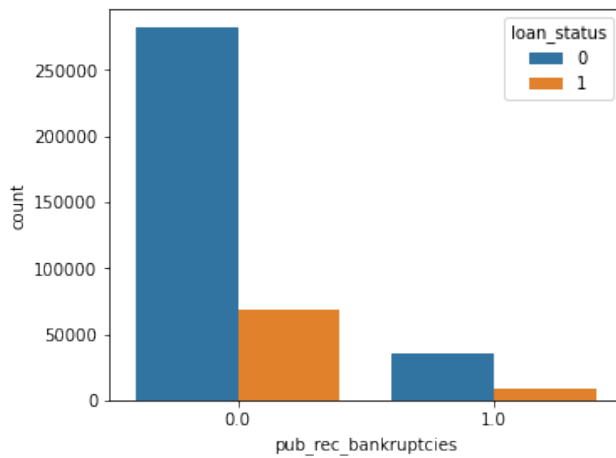

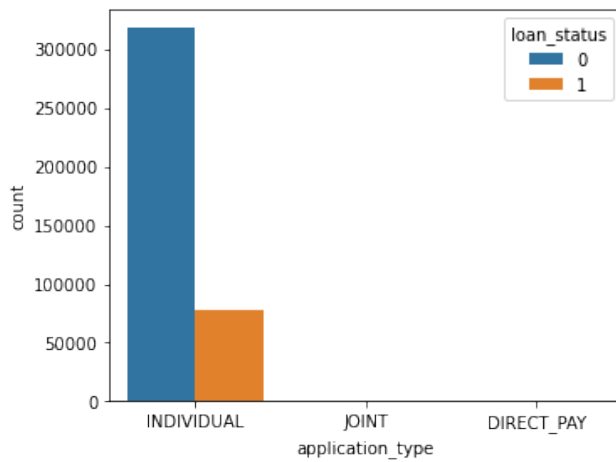


```

def pub_rec_bankruptcies(number):
    if number == 0.0:
        return 0
    elif number >= 1.0:
        return 1
    else:
        return number

data['pub_rec'] = data.pub_rec.apply(pub_rec)
data['mort_acc'] = data.mort_acc.apply(mort_acc)
data['pub_rec_bankruptcies'] =
data.pub_rec_bankruptcies.apply(pub_rec_bankruptcies)

plt.figure(figsize=(12, 30))
plt.subplot(6, 2, 1)
sns.countplot(x='pub_rec', data=data, hue='loan_status')
plt.subplot(6, 2, 2)
sns.countplot(x='initial_list_status', data=data, hue='loan_status')
plt.subplot(6, 2, 3)
sns.countplot(x='application_type', data=data, hue='loan_status')
plt.subplot(6, 2, 4)
sns.countplot(x='mort_acc', data=data, hue='loan_status')
plt.subplot(6, 2, 5)
sns.countplot(x='pub_rec_bankruptcies', data=data, hue='loan_status')
plt.show()

```



```
#lets impute the null value of mort_acc with the help of total_acc
data.groupby(by='total_acc')['mort_acc'].median()

total_acc
2.0    0.0
3.0    0.0
4.0    0.0
```

```
5.0      0.0
6.0      0.0
```

```
...
124.0    1.0
129.0    1.0
135.0    1.0
150.0    1.0
151.0    0.0
```

Name: mort_acc, Length: 118, dtype: float64

```
total_acc_avg = data.groupby(by='total_acc').median().mort_acc
```

```
total_acc_avg
```

```
total_acc
2.0      0.0
3.0      0.0
4.0      0.0
5.0      0.0
6.0      0.0
```

```
...
124.0    1.0
129.0    1.0
135.0    1.0
150.0    1.0
151.0    0.0
```

Name: mort_acc, Length: 118, dtype: float64

```
def fill_mort_acc(x):
    total_acc = x['total_acc']
    mort_acc = x['mort_acc']

    if np.isnan(mort_acc):
        return total_acc_avg[total_acc].round()
    else:
        return mort_acc
```

```
data['mort_acc'] = data.apply(fill_mort_acc, axis=1)
```

```
data.isnull().sum()
```

```
loan_amnt      0
term           0
int_rate       0
grade          0
sub_grade      0
emp_title      22927
emp_length     18301
home_ownership 0
annual_inc     0
verification_status 0
```

```
issue_d          0
loan_status      0
purpose          0
title            1755
dti              0
earliest_cr_line 0
open_acc         0
pub_rec          0
revol_bal        0
revol_util       276
total_acc        0
initial_list_status 0
application_type 0
mort_acc         0
pub_rec_bankruptcies 535
address          0
dtype: int64
```

```
data.shape
```

```
(396030, 26)
```

```
data.dropna(inplace=True)
```

```
data.shape
```

```
(370622, 26)
```

```
#Outlier Detection & Treatment -
```

```
numerical_data = data.select_dtypes(include='number')
```

```
num_cols = numerical_data.columns
```


```
len(num_cols)
```

```
12
```


```
def box_plot(col):
    plt.figure(figsize=(8, 5))
    sns.boxplot(x=data[col])
    plt.title('Boxplot')
    plt.show()
```

```
for col in num_cols:
    box_plot(col)
```


Boxplot




Boxplot




Boxplot




Boxplot




Boxplot




Boxplot




Boxplot




Boxplot




Boxplot




Boxplot



Boxplot





```
for col in num_cols:
    mean = data[col].mean()
    std = data[col].std()
    upper_limit = mean+3*std
    lower_limit = mean-3*std
    data = data[(data[col]<upper_limit) & (data[col]>lower_limit)]
data.shape
```

```
(354519, 26)
```

#Data Preprocessing -

```
term_values = {' 36 months': 36, ' 60 months': 60}
data['term'] = data.term.map(term_values)
list_status = {'w': 0, 'f': 1}
data['initial_list_status'] =
data.initial_list_status.map(list_status)
data['zip_code'] = data.address.apply(lambda x: x[-5:])
```

```
data['zip_code'].value_counts(normalize=True)*100
```

```
70466    14.382022
30723    14.277373
22690    14.268347
```

```

48052    14.127028
00813    11.610097
29597    11.537322
05113    11.516731
93700     2.774746
11650     2.772771
86630     2.733563
Name: zip_code, dtype: float64

```

#dropping the columns which are of less significance

```

data.drop(columns=['issue_d', 'emp_title', 'title', 'sub_grade',
                  'address', 'earliest_cr_line', 'emp_length'],
          axis=1, inplace=True)

```

```
data.shape
```

```
(354519, 20)
```

```
data.columns
```

```

Index(['loan_amnt', 'term', 'int_rate', 'grade', 'home_ownership',
       'annual_inc', 'verification_status', 'loan_status', 'purpose',
       'dti',
       'open_acc', 'pub_rec', 'revol_bal', 'revol_util', 'total_acc',
       'initial_list_status', 'application_type', 'mort_acc',
       'pub_rec_bankruptcies', 'zip_code'],
      dtype='object')

```

```
data
```

	loan_amnt	term	int_rate	grade	home_ownership	annual_inc	\
0	10000.0	36	11.44	B	RENT	117000.0	
1	8000.0	36	11.99	B	MORTGAGE	65000.0	
2	15600.0	36	10.49	B	RENT	43057.0	
3	7200.0	36	6.49	A	RENT	54000.0	
4	24375.0	60	17.27	C	MORTGAGE	55000.0	
...	
396025	10000.0	60	10.99	B	RENT	40000.0	
396026	21000.0	36	12.29	C	MORTGAGE	110000.0	
396027	5000.0	36	9.99	B	RENT	56500.0	
396028	21000.0	60	15.31	C	MORTGAGE	64000.0	
396029	2000.0	36	13.61	C	RENT	42996.0	

	verification_status	loan_status	purpose	dti
open_acc \				
0	Not Verified	0	vacation	26.24
16.0				
1	Not Verified	0	debt_consolidation	22.05
17.0				
2	Source Verified	0	credit_card	12.79
13.0				

3	Not Verified		0	credit_card	2.60
6.0					
4	Verified		1	credit_card	33.95
13.0					
...
...					
396025	Source	Verified	0	debt_consolidation	15.63
6.0					
396026	Source	Verified	0	debt_consolidation	21.45
6.0					
396027		Verified	0	debt_consolidation	17.56
15.0					
396028		Verified	0	debt_consolidation	15.88
9.0					
396029		Verified	0	debt_consolidation	8.32
3.0					
	pub_rec	revol_bal	revol_util	total_acc	initial_list_status
\					
0	0	36369.0	41.8	25.0	0
1	0	20131.0	53.3	27.0	1
2	0	11987.0	92.2	26.0	1
3	0	5472.0	21.5	13.0	1
4	0	24584.0	69.8	43.0	1
...
396025	0	1990.0	34.3	23.0	0
396026	0	43263.0	95.7	8.0	1
396027	0	32704.0	66.9	23.0	1
396028	0	15704.0	53.8	20.0	1
396029	0	4292.0	91.3	19.0	1
	application_type	mort_acc	pub_rec_bankruptcies	zip_code	
0	INDIVIDUAL	0.0	0.0	22690	
1	INDIVIDUAL	1.0	0.0	05113	
2	INDIVIDUAL	0.0	0.0	05113	
3	INDIVIDUAL	0.0	0.0	00813	
4	INDIVIDUAL	1.0	0.0	11650	
...	
396025	INDIVIDUAL	0.0	0.0	30723	

396026	INDIVIDUAL	1.0	0.0	05113
396027	INDIVIDUAL	0.0	0.0	70466
396028	INDIVIDUAL	1.0	0.0	29597
396029	INDIVIDUAL	1.0	0.0	48052

[354519 rows x 20 columns]

#--one hot encoding

```
columns_to_encode = ['purpose', 'zip_code', 'grade',
                      'verification_status', 'application_type', 'home_ownership']
```

Perform one-hot encoding

```
data = pd.get_dummies(data, columns=columns_to_encode,
                      drop_first=True)
data.shape
```

(354519, 49)

#Data Preparation for Modeling -

```
X = data.drop('loan_status', axis=1)
y = data['loan_status']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.30,
                                                    stratify=y, random_state=42)
```

```
scaler = MinMaxScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
logreg = LogisticRegression(max_iter=1000)
logreg.fit(X_train, y_train)
```

```
LogisticRegression(max_iter=1000)
```

```
y_pred = logreg.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print('Accuracy of Logistic Regression Classifier on test set:
{:.3f}'.format(accuracy))
```

Accuracy of Logistic Regression Classifier on test set: 0.890

```
confusion_matrix = confusion_matrix(y_test, y_pred)
print(confusion_matrix)
```

```
[[85365  523]
 [11131 9337]]
```

```
print(classification_report(y_test, y_pred))
```


	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.88	0.99	0.94	85888
1	0.95	0.46	0.62	20468
accuracy			0.89	106356
macro avg	0.92	0.73	0.78	106356
weighted avg	0.90	0.89	0.87	106356

```

logit_roc_auc = roc_auc_score(y_test, logreg.predict(X_test))
fpr, tpr, thresholds = roc_curve(y_test, logreg.predict_proba(X_test)[:,1])
plt.figure()
plt.plot(fpr, tpr, label='Logistic Regression (area = %0.2f)' %
logit_roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
plt.show()

```



```

def precision_recall_curve_plot(y_test, pred_proba_c1):
    precisions, recalls, thresholds = precision_recall_curve(y_test,
pred_proba_c1)

```

```


threshold_boundary = thresholds.shape[0]
# plot precision
plt.plot(thresholds, precisions[0:threshold_boundary],
linestyle='--', label='Precision')
# plot recall
plt.plot(thresholds, recalls[0:threshold_boundary],
label='recalls')

start, end = plt.xlim()
plt.xticks(np.round(np.arange(start, end, 0.1), 2))

plt.xlabel('Threshold Value'); plt.ylabel('Precision and Recall
Value')
plt.legend(); plt.grid()
plt.show()

precision_recall_curve_plot(y_test, logreg.predict_proba(X_test)[:,-1])

```



```

from statsmodels.stats.outliers_influence import
variance_inflation_factor
def calc_vif(X):
    # Calculating the VIF
    vif = pd.DataFrame()
    vif['Feature'] = X.columns
    vif['VIF'] = [variance_inflation_factor(X.values, i) for i in
range(X.shape[1])]
    vif['VIF'] = round(vif['VIF'], 2)

```



```

    vif = vif.sort_values(by='VIF', ascending = False)
    return vif
#calc_vif(X)[:5]

X.drop(columns=['application_type_INDIVIDUAL'], axis=1, inplace=True)
#calc_vif(X)[:5]

X.drop(columns=['int_rate'], axis=1, inplace=True)
#calc_vif(X)[:5]

X.drop(columns=['term'], axis=1, inplace=True)
#calc_vif(X)[:5]

X.drop(columns=['purpose_debt_consolidation'], axis=1, inplace=True)
#calc_vif(X)[:5]

X.drop(columns=['open_acc'], axis=1, inplace=True)
#calc_vif(X)[:5]

X = scaler.fit_transform(X)
kfold = KFold(n_splits=5)
accuracy = np.mean(cross_val_score(logreg, X, y, cv=kfold,
scoring='accuracy', n_jobs=-1))
print("Cross Validation accuracy: {:.3f}".format(accuracy))

Cross Validation accuracy: 0.891

from imblearn.over_sampling import SMOTE
sm = SMOTE(random_state=42)
X_train_res, y_train_res = sm.fit_resample(X_train, y_train.ravel())

print("After OverSampling, counts of label '1':
{}".format(sum(y_train_res == 1)))
print("After OverSampling, counts of label '0':
{}".format(sum(y_train_res == 0)))

After OverSampling, counts of label '1': 200405
After OverSampling, counts of label '0': 200405

lr1 = LogisticRegression(max_iter=1000)
lr1.fit(X_train_res, y_train_res)
predictions = lr1.predict(X_test)

# Classification Report
print(classification_report(y_test, predictions))

```

	precision	recall	f1-score	support
0	0.95	0.80	0.87	85888
1	0.49	0.81	0.61	20468
accuracy			0.80	106356

macro avg	0.72	0.80	0.74	106356
weighted avg	0.86	0.80	0.82	106356


```
def precision_recall_curve_plot(y_test, pred_proba_c1):
    precisions, recalls, thresholds = precision_recall_curve(y_test,
pred_proba_c1)

    threshold_boundary = thresholds.shape[0]
    # plot precision
    plt.plot(thresholds, precisions[0:threshold_boundary],
linestyle='--', label='precision')
    # plot recall
    plt.plot(thresholds, recalls[0:threshold_boundary],
label='recalls')

    start, end = plt.xlim()
    plt.xticks(np.round(np.arange(start, end, 0.1), 2))

    plt.xlabel('Threshold Value'); plt.ylabel('Precision and Recall
Value')
    plt.legend(); plt.grid()
    plt.show()

precision_recall_curve_plot(y_test, lr1.predict_proba(X_test)[: ,1])
```



#we should try our best to have improved precision recall both

High recall means that the model is effective at capturing all positive instances. This is important in scenarios where missing a positive instance has severe consequences, such as failing to detect a fraudulent transaction.

